Operating System Implementation

Thomas Hybel

Aarhus University
November 2017

Abstract

This report documents our implementation of an operating system from scratch, with the aim of furthering our understanding of operating system design and internals. The development process was separated into eight labs, each comprising a major area of functionality. In the first four labs we built the infrastructure necessary to run programs in user space. In the following labs we added a file system, networking support, and a graphical user interface to the operating system. We also moved the operating system from emulated hardware to a real machine.

CONTENTS

INTRODUCTION

This report documents our journey through writing an operating system from scratch. Our primary goal was to learn about operating system internals and design, and to investigate the feasibility of the task. Neither usability nor efficiency were significant goals throughout the process, except when their consideration facilitated learning.

To get started on kernel development we made extensive use of material from the MIT 2016 Operating System Engineering course, which is accessible online [1]. The MIT course is divided into six major "labs", with each lab encapsulating a set of kernel functionality. For example, lab 4 involves functionality needed for multitasking and concurrency, while lab 5 comprises the implementation of a custom file system.

The kernel we wrote is for the 32-bit x86 architecture. The kernel is written in C, with some parts having to be hand-written in assembly. We have used GCC as our compiler, Git for version control, and GDB for debugging. Additionally, it was convenient to develop the kernel on emulated hardware rather than a real machine, since this eased debugging and increased development speed. It also made it trivial to change the amount of system RAM, and to add new processors or an extra hard drive. We have used the full-system emulator QEMU for this task.

Each MIT lab has an associated web page which describes what needs to be done and provides links to resources like manuals and specifications. Since we opted to follow the MIT course, this also meant that some design decisions were made for us. One example of this is us following an exokernel philosophy, which meant that we have pushed large parts of the kernel functionality into user space.

During kernel development, occasionally some functionality was important, but not technically interesting to implement. In these cases the MIT lab often provided the code for us, with enough code missing that we had to understand the concept to complete it, while still saving significant amounts of time. An example of this is the code for communicating with the outside world over a serial connection, which is used throughout the labs.

In addition to the six labs of the MIT course, we have decided upon two labs on our own. In lab 7 we designed and implemented a graphical user interface for the operating system. In lab 8 we made the operating system run on real hardware instead of QEMU. Since we were given zero guidance nor code, these labs were significantly more difficult and time-consuming than previous ones.

Each of the following sections corresponds directly to one lab. They describe the development process, from first booting into a minimal kernel, until the end when we run a full graphical operating system on a real machine.

BOOTING

In this lab we wrote initialization code and a boot loader for our kernel. The initialization code switches the processor to 32-bit protected mode. The boot loader loads the rest of the kernel into memory and jumps to it.

2.1 THE BOOT PROCESS

To understand the purpose of a boot loader, it helps to have an overview of the startup process of an x86 machine. When an x86 machine starts, its BIOS code runs. The BIOS initializes some of the hardware, and then it loads the first sector (512 bytes) from the boot medium into a hard-coded address which it jumps to.

This first sector will typically contain a small program known as the boot loader. Its purpose is to load the main kernel from disk and transfer execution to it.

Before the boot loader loads the kernel, the kernel should first run some initialization code which sets up a more comfortable environment for the boot loader and kernel to work in.

2.2 INITIALIZATION CODE

When the BIOS jumps into our code, the processor is running in 16-bit real mode. However the code produced by a modern compiler expects to run in 32-bit protected mode. The initialization code should therefore be written in assembly and should switch processor modes.

The most important difference between real and protected mode is that real mode does not support the use of a page table to implement virtual memory.

We switch to protected mode by setting a bit in the control register `cr0`. Enabling virtual memory similarly works by setting another bit in `cr0`. We do not immediately enable virtual memory, though, since we do not yet have the infrastructure to set up a proper page table. This happens in section 3. Until then, all addressing is physical.

To switch to 32-bit mode, we must update the `cs` (code segment) register. The `cs` register is an offset into a table called the Global Descriptor Table (GDT) which is an array of descriptors [2]. Each descriptor describes a segment of memory; a bit in the descriptor determines whether the segment contains 16-bit or 32-bit code.

We use the `ljmp` instruction to update the `cs` register and use a 32-bit segment descriptor. Next, the boot loader is executed; since the processor is in 32-bit protected mode, the rest of the kernel code can be written in C rather than assembly.

2.3 THE BOOT LOADER

It is the task of the boot loader to load the main kernel and transfer execution to it. Since the boot loader resides on the first sector of the hard drive, it must load the kernel using no more than 512 bytes of code, minus the bytes used by the initialization code.

Our kernel is compiled into an Execute and Linkable Format (ELF) file. The ELF file specifies the physical addresses into which the code and data of the kernel must be loaded. The boot loader must thus parse the ELF file to load the kernel.

Once the boot loader has loaded the kernel, it determines the entry point address from the ELF file and jumps there.

2.4 MINIMAL KERNEL CODE

At this point we were able to run kernel code. However we did not yet have any functionality; unlike user-mode programs, the kernel does not have access to a C standard library, unless we write one ourselves.

As a start, we wanted the ability to input and output text. Our kernel uses the `inb` and `outb` instructions to communicate with the outside world via a serial connection. We used this to print a "hello, world" message and confirm that the kernel runs.

Using instructions such as `inb` and `outb` to communicate with an input/output device is quite common. The method is known as Programmed Input/Output (PIO); it will be used often in the following sections.

MEMORY MANAGEMENT

The goal of this lab was to enable virtual memory after setting up a page table. To set up the page table, we first needed to implement a subsystem for allocating and freeing pages of physical memory.

3.1 PHYSICAL PAGE MANAGEMENT

A given system has a limited amount of physical memory, depending on how much RAM the machine has. This memory is split up into a number of pages. On x86 a page is 4096 bytes, or 0x1000 in hexadecimal. Thus pages always aligned on 0x1000-byte boundaries. The kernel determines the amount of RAM by using PIO to query a memory area called the CMOS [3].

The physical page management subsystem will keep a reference count for each page. If the count is zero, the page is free and can be allocated. This metadata is stored in an array of `PageInfo` structs. Each entry in this array directly corresponds to one physical page of memory, such that the first `PageInfo` struct holds metadata about the first page of physical memory, and so on.

The `PageInfo` of free pages are additionally stored in a linked list, such that the kernel can return a free page in constant time.

We wrote the following functions to manage physical pages:

- `page_alloc` is used to allocate a page of physical memory

- `page_free` is used to put a page on the free list

- `page_decref` and `page_incref` are used to manage reference counts of pages

These functions provide critical infrastructure needed by other kernel features.

3.2 PAGE TABLE THEORY

It is necessary to introduce some theory before we can explain how our kernel initializes its page table.

The x86 page table is a two-level table whose main purpose is to let the processor translate a virtual address to a physical address [4]. The page table can be thought of as a 1024-ary tree with two levels.

A pointer to the first level of the page table can be found in the `cr3` register. The first level is called the Page Directory. It contains 1024 Page Directory Entries (PDEs). Each PDE points to a second-level node, which contains 1024 Page Table Entries (PTEs). A PTE specifies a page of physical memory and its permissions, including whether it is writable and whether it is accessible to user-mode code.

To translate from a virtual to a physical address, it is necessary to walk the page table. Say that the process wishes to access a virtual address v. It first looks in the `cr3` register to find the Page Directory. It uses the 10 higher-order bits of v to specify a PDE. It uses the next 10 bits of v to specify a PTE The PTE contains the address of a physical page. The last 12 bits of v are used as an offset into this page, and the address translation is complete. The processor also validates the permissions of the page before the access, and generates a page fault if these are inappropriate.

This is a costly process, and in practice the job is done by specialized hardware called a Memory Management Unit (MMU). Additionally, recent translations are cached in the so-called Translation Lookaside Buffer (TLB).
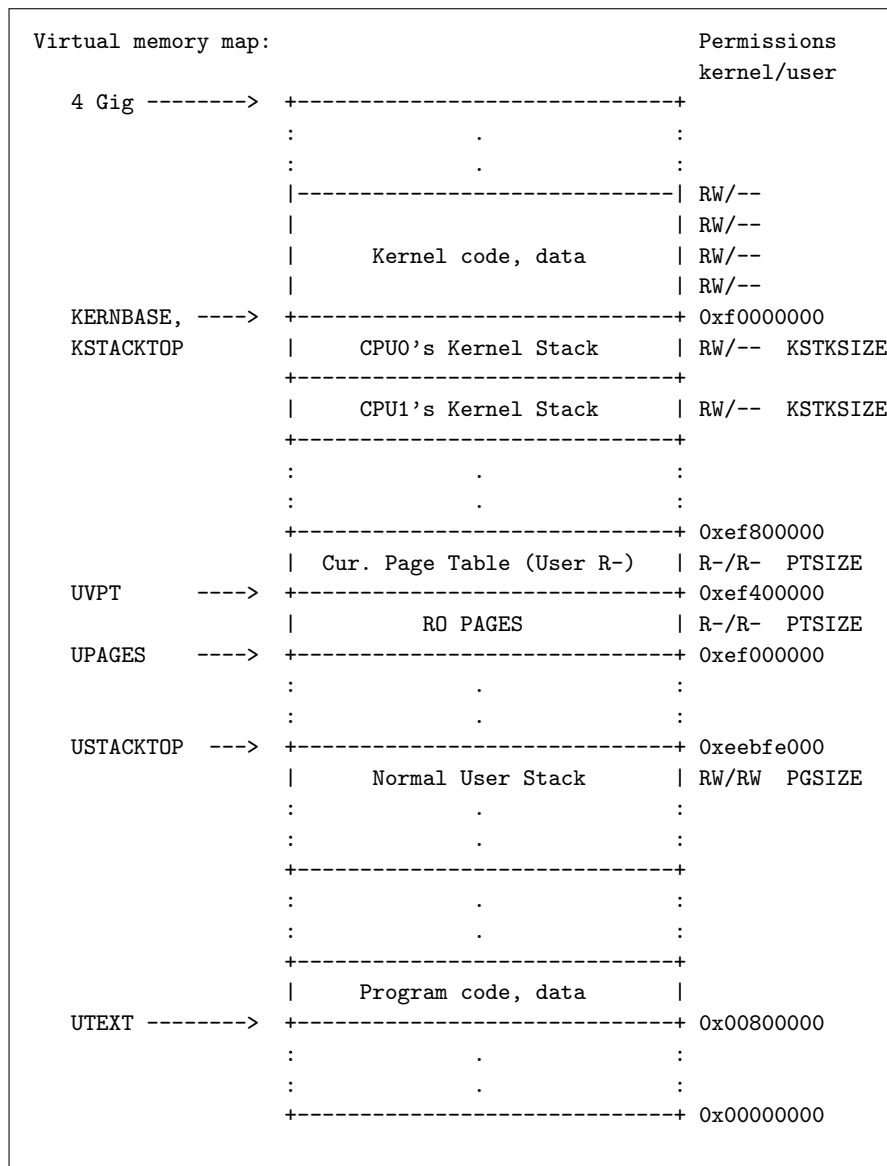
```
 Virtual memory map:                                Permissions
                                                    kernel/user
    4 Gig -------->  +----------------------------+
                     :              .             :
                     :              .             :
                     |----------------------------| RW/--
                     |                            | RW/--
                     |       Kernel code, data    | RW/--
                     |                            | RW/--
    KERNBASE, ---->  +----------------------------+ 0xf0000000
    KSTACKTOP        |     CPU0's Kernel Stack    | RW/--   KSTKSIZE
                     +----------------------------+
                     |     CPU1's Kernel Stack    | RW/--   KSTKSIZE
                     +----------------------------+
                     :              .             :
                     :              .             :
                     +----------------------------+ 0xef800000
                     | Cur. Page Table (User R-)  | R-/R-   PTSIZE
    UVPT      ---->  +----------------------------+ 0xef400000
                     |            RO PAGES        | R-/R-   PTSIZE
    UPAGES    ---->  +----------------------------+ 0xef000000
                     :              .             :
                     :              .             :
    USTACKTOP --->   +----------------------------+ 0xeebfe000
                     |        Normal User Stack   | RW/RW   PGSIZE
                     :              .             :
                     :              .             :
                     +----------------------------+
                     :              .             :
                     :              .             :
                     +----------------------------+
                     |      Program code, data    |
    UTEXT -------->  +----------------------------+ 0x00800000
                     :              .             :
                     :              .             :
                     +----------------------------+ 0x00000000
```

Figure 3.1: The virtual address space of the kernel

3.3 PAGE TABLE MANAGEMENT

We wrote the following functions to manage the page table:

- `pgdir_walk` is called by most of the other functions to walk the page table, finding the PTE corresponding to a given virtual address. It allocates new levels of the page table as needed using `page_alloc` from the previous section.

- `page_insert` is used to insert a physical page into a page table at a given virtual address. In other words, it finds a PTE using `pgdir_walk` and stores the physical address there.

- `page_lookup` finds the physical address of a page, given a virtual address.

- `page_remove` invalidates a PTE in a page table.

To set up the page table, the kernel allocates pages of physical memory and uses the newly implemented functions to insert these into the table.

The memory layout of the address space of the kernel is largely up to us. Figure 3.1 gives a simplified overview of the layout we opted for. The diagram

shows that the kernel code resides starting at virtual address `0xf0000000`. The kernel stacks, used by processors when running kernel-mode code, reside just below, between `0xefc00000` and `0xf0000000`. Further down in the address space, between `0xef400000` and `0xef800000`, we have the User Virtual Page Table (UVPT) area, which gives user-mode processes read-only access to their page table, enabling certain exokernel-style programs to work. The stack of the user-mode program starts at `0xeebfe000` and grows towards lower addresses. The user-mode code and data reside near the bottom of the address space, around `0x00800000`.

We created a page table according to this layout. We then updated the `cr3` register and set a bit in the `cr0` register to enable the use of the page table. At this point our kernel had proper virtual memory. Once user-mode processes have been implemented, virtual memory guarantees that the processes cannot modify the address space of one another, nor can they corrupt the kernel.

4

The goal of this lab was to run a user-mode process. This required us to write infrastructure for managing processes and their metadata. We also modified our kernel to handle any exceptions generated by user-mode processes. Finally we implemented a system call mechanism to let processes interact with the kernel.

4.1 MANAGING PROCESS METADATA

Each process has some associated information. This includes its state (running, runnable, blocked, killed), its process ID, parent process ID, page directory, saved register state, and so on. This metadata is stored in a struct `Env`. The process subsystem is similar to the physical page subsystem; an array holds the `Env` struct of each process on the system, and free environments are stored in a linked list. We implemented functions for creating, initializing, and destroying a process.

4.2 ELF LOADING

Programs are represented as ELF files. To launch a process, the kernel first allocates a fresh page table. It then walks over each section in the ELF file, reads at which virtual address the section should go, allocates corresponding physical pages, inserts them into the page table, and copies the code or data into the physical pages.

Note that we have not yet introduced a file system, so it is not immediately clear where the kernel can find the programs which it should load. To solve this problem we embed each user-mode program into the kernel as a blob of binary data. In section 6 we describe our implementation of a proper file system for storing programs and data.

4.3 CONTEXT SWITCHING

Once a program has been loaded, the kernel must perform a context switch to let the created process run. During a context switch the kernel restores the saved general-purpose registers using the `popal` instruction. It also updates `cr3` to switch to the new address space. Finally it uses the `iret` instruction to restore the saved `eip`, `esp`, `eflags`, and `cs` registers. This transfers execution to user-mode code.

The kernel must also drop its privileges. The lower two bits of the `cs` register determine the privilege level of the processor. This was previously 0, since the kernel runs in ring 0. By setting this to 3 during the `iret`, the processor switches to user-mode operation, which is ring 3. Since instructions such as `iret` may only be run from privilege level 0, the process cannot simply modify its `cs` register to increase its privileges.

At this point our kernel was able to successfully create a new process, load program code and data into its address space, and let the process run. Unfortunately the code had no way to give control back to the kernel, so it simply ran forever, or at least until it triggered an exception or interrupt. Since this was not yet handled, any exception caused the whole system to crash.

4.4   HANDLING OF EXCEPTIONS AND INTERRUPTS

The processor may trigger an exception while processing an instruction, e.g., on a division by zero or illegal memory access [5, 6]. The processor may also occasionally trigger an interrupt. This often happens for asynchronous reasons, such as when a key is pressed or a network packet is received. Interrupts can also be raised with the `int` instruction; this mechanism can be used to implement system calls.

The result of an exception or interrupt is that the processor enters kernel mode through a context switch, running exception- or interrupt-specific handler code. This handler code is found as follows. Exceptions and interrupts are numbered. This number specifies an index into a table called the Interrupt Descriptor Table (IDT), which can be found using the IDT register, `idtr`. The IDT entry describes the address of the handler code.

We wrote a common function, `trap`, which is called whenever an exception or interrupt occurs. For each exception and interrupt, we filled in its IDT entry with a small stub which passes the number of the exception or interrupt as the first argument and calls `trap`.

The typical result of an exception is that the kernel terminates the running process. If an exception occurs in kernel mode, the result is a kernel panic, where the kernel prints an error message and hangs.

4.5   HANDLING OF SYSTEM CALLS

A process frequently needs to call kernel code, e.g. to print text onto the screen or to perform inter-process communication. This is called a system call.

One mechanism to implement system calls on the x86 architecture is to use the `int` instruction, which triggers an interrupt when executed. We are free to use any interrupt that is not in use, so we arbitrarily chose number `0x30`. Thus a program uses the `int 0x30` instruction to perform a system call. We chose a fairly standard ABI; the `eax` register holds the system call number, while arguments go in registers `ebx`, `ecx`, etc.

The `trap` function recognizes interrupt number `0x30` and calls the `syscall` function, which uses a large `switch` to delegate each system call to a specific handler.

We wrote kernel system call handlers for input and output of a single character, as well as for process termination. User-mode programs did not have a C standard library at this point, so we created one, adding an interface to the system calls to the library. At this point we were able to run a simple user-mode program and have it interact with the user.

MULTIPROCESSING

In this lab we added the features necessary for running multiple processes concurrently and having them interact.

We implemented a simply round-robin scheduler. The scheduler can preempt the running process when its time slice is up with the aid of a device called a LAPIC, which can generate periodic timer interrupts.

We also wrote code to activate any processors beyond the first, letting processes run truly concurrently.

We implemented a naive fork mechanism which lets proceses spawn more processes. We then upgraded its efficiency with a copy-on-write mechanism.

Finally, we added an inter-process communication feature, letting processes send values and pages of memory to one another.

5.1   THE PROCESS SCHEDULER

The main purpose of a scheduler is to decide which process to schedule in when the running process is scheduled out. There are many ways to make this decision, and the scheduler has a great impact on system responsivity and efficiency.

However for simplicity reasons, we opted for a simple round-robin scheduler. That is, the scheduler keeps a circular queue of all processes. The first runnable process in the queue is chosen to be scheduled in.

The scheduler also needs to preempt each process when its time slice is used. To accomplish this, during kernel initialization the kernel asks a device called the LAPIC to raise a timer interrupt periodically, waiting some fixed amount of bus cycles between each raised interrupt. Our `trap` function recognizes this timer interrupt and reacts by asking the scheduler to schedule in a new process. Interaction with the LAPIC is described in section 5.4.

5.2   ACTIVATING MORE PROCESSORS

So far the kernel has run on an emulated single-core machine. However a machine with n processors can be emulated by passing the `-smp n` option to QEMU.

When a system with multiple processors boots, the hardware dynamically selects only a single processor to run [7, section 8.4]. It is called the Bootstrap Processor (BP) while the remaining processors, if any, are called Application Processors (APs). It is up the the BP to start up the APs when the system is ready.

We wrote code which starts up the APs. This is accomplished by querying the LAPIC of the BP, having it send an inter-processor interrupt to each AP. Upon receiving such an interrupt, the APs start executing code at an address specified by the AP.

The APs start in 16-bit real mode, just as the BP did. They therefore need to switch to 32-bit protected mode mode. After doing so, each AP calls into the scheduler to run a new process.

5.3   ENSURING MUTUAL EXCLUSION

With multiple processors running concurrently, all the typical issues of concurrency arose. Multiple processors could modify kernel data simultaneously, leading to race conditions.

We prevented this with a trivial but inefficient approach: we added a lock which must be locked before running any kernel code. Thus only one

processor may run kernel code at a time. Still, user-mode processes can run truly concurrently.

The kernel lock is a spinlock. It is a global variable. A processor repeatedly uses the `lock` and `xchg` instructions to atomically exchange the global variable with the value 1. If the old value was zero, the processor now holds the lock and may enter the kernel. Otherwise it retries.

We added calls to lock and unlock the kernel in the right places. At this point our kernel was capable of running multiple user-mode processes concurrently.

5.4   INTERACTING WITH THE LAPIC

In section 5.1 the scheduler had to ask the LAPIC to generate timer interrupts for it, and in section 5.2 the BP used the LAPIC to send inter-processor interrupts to the APs. We therefore needed code to query the LAPIC.

First, however, we explain the acronym. A Programmable Interrupt Controller (PIC) is a hardware device responsible for managing interrupts for the processor. For example, if multiple interrupts are generated simultaneously, the PIC can prioritize the interrupts and deliver them one at a time. When Intel updated their PIC standard to include new features, the conforming device was called an Advanced PIC (APIC). The APIC has a component called the Local APIC (LAPIC) which is local to each processor.

The processor communicates with the LAPIC using Memory-Mapped I/O (MMIO) [8]. This means that the LAPIC is mapped into memory at a specific, system-dependent physical address. Reading at certain offsets will correspond to reading from certain registers in the LAPIC and likewise for writing. Thus to communicate with the LAPIC, the kernel merely needs to read and write to certain addresses. The difficult part is figuring out the physical address where the LAPIC resides.

There are multiple ways to find the LAPIC. For now, our kernel uses the method described in the multi-processor specification [9], which we will refer to as the mpconfig method. It involves searching through parts of physical memory to find a structure called the MP floating pointer structure. This structure points to a table called the MP configuration table, which contains the physical address of the LAPIC.

5.5   A SIMPLE FORK MECHANISM

So far, every process was directly spawned by the kernel. However, a process should also be able to spawn processes. We therefore needed a mechanism to let a process fork. In this section we describe our initial, simple implementation of fork, and in the following section we improve it by introducing a copy-on-write mechanism.

Since our kernel is an exokernel, we prefer to keep code outside of kernel land. We therefore wrote a number of system calls which can be combined to implement a user-mode fork. Specifically, we wrote the following system calls:

- `sys_exofork` creates a non-runnable child process with an empty address space.

- `sys_env_set_status` can mark a process as runnable.

- `sys_page_alloc` allocates an empty page in the address space of a process.[1]

- `sys_page_map` maps a page from the current process into a child process.

---

[1] For security reasons, the system calls are coded to ensure that a process cannot modify pages in unrelated, non-child processes.

- `sys_page_unmap` unmaps a page from the current process or a child process.

Besides these system calls, it is also necessary for a process to have access to information about the layout of its own address space. This is already the case; in section 3.3 we set up the page table of a process such that part of the address space contains its page table.

To fork, a parent process goes through the following steps:

- The parent process calls `sys_exofork` to create a new child process with an empty address space which is not initially runnable.

- The parent walks over its page table, and for each mapped page, it does the following:

  - The parent uses `sys_page_map` to create a temporary page at a temporary address.

  - The parent copies the contents of the current page into the temporary page.

  - The parent uses `sys_page_map` to insert the temporary page into the address space of the child process at the original address.

  - The parent uses `sys_page_unmap` to remove the temporary page from its own address space.

- The parent marks the child as runnable using `sys_env_set_status`.

At this point the fork is complete, and since the child is marked as runnable, it will eventually be scheduled in.

5.6 COPY-ON-WRITE FORK

The fork mechanism described in the previous section is slow and memory-inefficient, because it indiscriminately copies every page of the parent into the child process. This involves a lot of reading and writing of RAM, and it means that if the parent used n physical pages of memory, then after a fork, $2n$ physical pages will be used.

However the same physical page can transparently be mapped into both the parent and child process, as long as it is never modified. In fact, *all* the pages in the child process can initially be shared with the parent. It is only once a write happens that a page must be copied. We have implemented such a copy-on-write fork mechanism almost entirely in user land, following the exokernel design philosophy.

To perform as much work in user land as possible, we implemented a mechanism which lets a process handle its own page faults. By default a page fault will result in process termination. However we have implemented a system call, `sys_env_set_pgfault_upcall`, which lets a process set a handler function. Then the kernel will handle a page fault by modifying the saved `eip` and `esp` registers of the process, pushing the old register values onto an exception stack, and switching the process back in.

When a process forks, *all* pages are initially shared between the parent and child process. However writable pages have their writable bit removed from their page table entry (PTE). Instead we set another bit which marks the page as copy-on-write.

When the child or parent process attempts to write to one of the now-shared pages, a fault will occur since the page is not writable anymore. The kernel sees that the copy-on-write bit is set, and then it delegates to the registered user-mode handler. The handler then uses the same method as in the simple fork implementation to map a new writable page, copy the contents of the old page onto it, and replace the copy-on-write page with the new writable page.

5.7 INTER-PROCESS COMMUNICATION

Since this project follows the exokernel philosophy, many future features will reside in user land. Two examples are a file system daemon and a daemon implementing a network stack. Other processes need an inter-process communication (IPC) mechanism to make use of these daemons.

We therefore implemented two system calls, `sys_ipc_recv` and `sys_ipc_try_send`. When a process calls `sys_ipc_recv` it will hang, waiting to receive data. `sys_ipc_try_send` will send data to a process in a non-blocking fashion. By default, a 32-bit integer is sent between processes, but for efficiency an extra argument to the system calls allows the sender to share a full page of memory per system call.

This marks the end of the multiprocessing lab. Our kernel can now run user-mode processes in a truly concurrent fashion. A scheduler manages the running processes, preempting them when necessary, and processes can efficiently fork and communicate via IPC.

FILE SYSTEM

The goal of this lab was to implement a custom file system and a user-mode shell program.

6.1 FILE SYSTEM DESIGN

Having a file system lets programs store data persistently on disk. It also provides a storage place for programs, instead of embedding them directly in the kernel as we have done so far. In this section we describe our design of a custom file system.[1]

A file system can be thought of as a way to manage how files, folders, and their metadata are stored on a disk. A raw disk is essentially a portion of memory which can be read from and written to. It is customary to partition the memory of a disk into fixed-size blocks. For our file system, the block size will be 4096 bytes.

One of these blocks is special, since it contains metadata for the file system. This block is called the superblock. This metadata includes such things as the disk size and where to find the root folder.

Our file system is laid out as follows. Block 0, is not used by our FS; it is reserved for the boot loader. Block 1 is the superblock. The next few blocks, starting at block 2, hold a bitmap which determines whether the remaining blocks on disk are in use or free. The remaining blocks are used to store the concrete files and folders.

A file is represented as a `struct file`, which is stored in its own block. Such a `file` struct contains metadata, such as the file name and size. The struct also has 10 pointers to blocks that hold the raw file data. If the data cannot fit in 10 blocks, the `file` struct has a pointer to a block which holds another 1024 pointers to data blocks. Thus our file system has a maximum file size of $(10 + 1024) * 4096 = 4235264$ bytes, i.e., around 4 MB.

A folder is represented as a `struct folder`, which is exactly identical to a `struct file`, except that the $10 + 1024$ block pointers no longer point to raw data, but to other blocks holding `file` or `folder` structs. There is a type flag which allows distinction between files and folders.

We used a small script to create a raw disk image with an initialized file system of the described format. The script let us add files, such as sample programs, to the file system. We then attached this raw disk to QEMU.

6.2 FILE SYSTEM DAEMON

In accord with exokernel design, we let all file system interaction go through a user-mode process which we call the file system daemon.

The daemon interacts with the disk using PIO. However a normal user-mode process cannot use the `inb` and `outb` instructions to perform PIO, so our kernel needs to give the daemon I/O privileges. It does so by setting a bit in the `eflags` register while spawning the daemon.

The disk knows nothing of the file system which is stored upon it. The daemon can merely read a sector of the disk at a time. It is therefore up to the daemon to implement reading and writing of files and folders according to the file system specification. We wrote the necessary functions for interacting with the file system.

The file system daemon spends its time looping, waiting for other processes to contact it via IPC. Processes can send requests to open, read, write,

---

[1] To keep our file system simple, we have left out many potential features, such as file and folder permissions, symbolic and hard links, and timestamps.

and stat files. The IPC details are hidden inside the user-mode C standard library, giving processes the usual interface with functions such as `open`, `read` and `write`.

For sake of illustration, the following happens when a process wants to open a file:

- The process calls the `open` library function.

- The library uses IPC to contact the file system daemon.

- The daemon reads the superblock to find the block of the root folder.

- The daemon follows pointers from folder to folder according to the given path.

- When the path has been traversed, the daemon has found the block that contains the `file` struct.

- The offset into the disk that holds the `file` struct is stored in a file descriptor.

- File descriptors are numbered; the daemon returns this number as the result of the IPC call.

When the process subsequently reads from the file descriptor, the daemon uses the pointers in the `file` struct to find the raw data and return it via IPC.

We have also implemented a block caching system to improve the efficiency of the file system daemon. When a block is first read, its contents are stored in RAM, and subsequent reads do not need to interact with the disk until the cache entry is invalidated through a write.

6.3 SHELL

To test the interaction between programs and the disk, we implemented a simplistic shell which allows loading and execution of programs from disk. The shell allows reads from and writes to the disk through input redirection (<) and output redirection (>). The shell additionally supports piping of the output of one program into another program. A small program, "cat", lets the user read files from the disk. A sample interaction is shown in figure 6.1.

```
$ ls
cat
echo
ls
sh
$ echo "Hello, OS." > motd
$ cat motd
"Hello, OS."
```

Figure 6.1: A sample interaction with the shell

NETWORKING

The goal of this lab was to connect our kernel to the internet by writing a network card driver.

7.1 THE NETWORK CARD

The operating system connects to a network using a network card, whose task it is to send and receive packets by interacting with the physical layer. QEMU emulates the Intel e1000 network card, which we therefore targeted. We added the card to our emulated machine with the "-net nic,model=e1000" option to QEMU, which creates a virtual router at IP 10.0.2.2 and assigns the guest an IP of 10.0.2.15. This let us hardcode the the IP address rather than having to implement a DHCP client or similar.

The e1000 manual [10] describes in detail the internals of the card, as well as the steps a driver must take to initialize the card, read incoming packets, and transmit outgoing packets.

The e1000 maintains two circular packet queues: a transmit queue and a receive queue. To send a packet, the kernel adds it to the transmit queue, which the network card periodically drains. When a packet arrives on the wire, the network card adds it to the receive queue, which the kernel periodically drains.

The items in the queues are not raw packets, but rather small descriptor structs, each of which describes an area of physical memory which can hold a packet. The queues are implemented as arrays, with each queue having a head and tail register to keep track of the head and tail of the queue.

7.2 THE NETWORK CARD DRIVER

We wrote an e1000 network card driver using the information from the manual. Our kernel uses PIO to scan the PCI bus, iterating over each connected device until it finds one whose vendor and device ID indicate it being an e1000 network card. The PCI interface supplies a physical address where the driver can use MMIO to communicate with the e1000.

With communication possible, the first task of the driver was to initialize the network card. The driver allocates the arrays that make up the packet queues and fills these with valid descriptors. It zeroes out the head and tail registers.

Once our driver has initialized the card, it writes into a register to enable it. It is then possible to transmit and receive packets. Our driver implements the code for taking packets from the receive queue and inserting packets into the transmit queue. User-mode programs have access to this driver functionality through two system calls, sys_receive and sys_transmit, respectively.

If the transmit queue ever becomes full due to the card draining it too slowly, it is up to the driver what should happen. For simplicity we have chosen to simply drop the packet if the queue is full. This is not a problem because the higher-level protocols are resistant to packet loss. Alternatively the driver could have blocked, waiting for space in the queue. If, on the other hand, the receive queue is full, the e1000 is designed to simply drop further packets.

7.3 THE NETWORK DAEMON

With the driver written, our kernel was capable of transmitting and receiving packets. However most processes only know of the data they wish to send; they are not capable of constructing packets.

We therefore needed a TCP/IP stack. Allegedly, writing such a network stack from scratch is an immense task. We therefore used the open-source stack lwIP ("lightweight IP"). lwIP acts as a black box for our purposes, taking raw data as input, and producing packets as output.

Rather than adding lwIP to the kernel, we embedded it in a network daemon. The daemon is responsible for managing sockets, just as the file system daemon was responsible for managing file descriptors. The network daemon takes in requests to send data via IPC, produces packets, and hands these over to the operating system via the `sys_transmit` system call. Likewise, it takes in requests to receive data and uses `sys_receive`, parses the resulting packets, and hands over the data to the requesting process.

The IPC communication is hidden away in the C standard library, such that user-mode programs have access to the familiar `connect`, `send` and `recv` interface.

7.4 WEB SERVER

To test the new functionality, we wrote a simple web server which can serve files from the file system over HTTP. The web server uses the network daemon to accept incoming connections and receive HTTP requests. It parses each request, queries the file system daemon to retrieve the requested file contents, and uses the network daemon to send back the HTTP response.

We configured QEMU to forward requests at port 80 to the emulated machine. At this point we were able to point a browser at the machine and be served a working web page.

GRAPHICS

In this lab we implemented a graphical user interface to our operating system. The end result is that each application is given its own graphical window on the screen, and the user can interact with applications using the mouse and keyboard.

The pixels on the screen are represented in a data structure called the Linear Frame Buffer (LFB), which the kernel must get access to by querying the BIOS. The kernel can draw pixels on the screen by writing to the LFB. We wrote a graphics library to draw pixels and more complex shapes.

We then designed a graphics stack, which involved deciding how input and output events should flow through the system from kernel to user application. After reading about how graphics work in other systems [11, 12, 13] we decided upon a design with a central display server, which is responsible for rendering all of the screen and which acts as an intermediary for input and output events.

Finally we added a few sample graphical applications to demonstrate the new features.

8.1 MAPPING THE LINEAR FRAME BUFFER

To initialize graphics, the operating system must query the BIOS to select a video mode, which is a description of the width, height and depth of the screen. Once a video mode is selected, the BIOS maps a crucial data structure into memory, which is called the Linear Frame Buffer (LFB).

The LFB represents the pixels of the screen; it can be thought of as a two-dimensional array, where each value is a 32-bit integer representing the RGB value of a pixel on the screen. Thus to write a pixel to the screen, the kernel must simply calculate the correct offset into the LFB and write a 32-bit integer there.

The kernel queries the BIOS to set a video mode using the `int 0x10` instruction. This transfers execution to the BIOS. However since the BIOS is written as 16-bit real mode code, our kernel must switch the processor back to 16-bit real mode before it can issue the interrupt. We therefore wrote assembly code which switches the processor mode, queries the BIOS to enumerate the valid video modes, and selects one with a satisfactory resolution.

At this point our kernel was able to draw pixels on the screen by writing to the LFB. However the kernel lacked the ability to draw more complex shapes.

8.2 GRAPHICS LIBRARY

We used the ability to draw a single pixel as a building block, writing functions that can draw more complex shapes, such as straight lines, rectangles, and windows. We also implemented font rendering.

Since this functionality is needed by many different applications, we put it inside a graphics library which is available to user space applications.

8.3 GRAPHICS STACK DESIGN

To implement a user interface, we needed to decide how input and output events should flow through the system. When a mouse click or key press occurs, the kernel receives an interrupt. How should this event arrive at a user space application, and which should it be? Should a user space

application render its own pixels? If so, how does it know which parts of the LFB it can write to without clashing with other applications?

After considering these questions and reading about how other operating systems render graphics, we decided to let all events flow through a central privileged process, called the display server.

The display server takes input events from the kernel and delivers them to the target application. Each application has a canvas. The display server takes output events from the applications, in the form of changes to the canvas, and writes the changes to the LFB.

Each graphical application waits to receive events from the display server in an event loop. The application can make changes to its canvas on the basis of an event occurring.

The kernel keeps a queue for input events. When raw input packets arrive from the mouse or keyboard, a driver handles these packets by putting them into an event queue. The display server periodically drains this queue through a system call, forwarding the events.

Figure **??** shows an overview of the different components in the graphics stack and how they interact.

The following sequence of events exemplifies the flow of events through the graphics stack:

- The user presses the 'A' key.

- The keyboard generates an interrupt to signal to the kernel that input is available.

- The kernel keyboard driver reads the pressed key using PIO.

- The driver puts the event into the events queue.

- The display server eventually drains the events queue and finds the key press event.

- The display server finds that the terminal application is active and therefore forwards the event to it using IPC.

- The terminal application receives the event and adds the 'A' to a buffer which holds the current input of the user. This buffer will eventually be used to run a command once the user presses the enter key.

- The terminal application also wants to display the 'A' on the screen, so that the user can see what is written so far. The application therefore asks the graphics library to render an 'A' on its canvas using the default font.

- The display server writes the pixels from the canvas into the part of the LFB that contains the application window.

- The user sees the 'A' appear on screen.

We describe the details of how we have implemented the display server in the following section.

8.4   THE DISPLAY SERVER

Various events must be performed periodically by the display server: it must drain input events from the kernel and deliver these to applications, and it must render the canvases of applications by writing to the LFB. We show the code for the main loop of the display server in figure 8.1.

The display server is responsible for writing the pixels of every application into the LFB. For efficiency the display server bypasses the kernel and writes directly to the LFB, which is accomplished by mapping the LFB into the address space of the display server using a new system call, sys_map_lfb.

```
while (1) {
        // drain input events from the kernel
        // and decide how to handle them
        process_events();

        // deliver events to targeted applications
        transmit_events();

        // draw the desktop background
        draw_background();

        // draw each application in the correct order
        draw_applications();

        // finally, draw the cursor on top
        draw_cursor();

        // write the changes to the LFB
        refresh_screen();
}
```

Figure 8.1: The main loop of the display server

The display server is responsible for spawning every other graphical application and assigning the application a portion of the screen. When spawning an application, an area of shared memory is set up through IPC. This memory is the canvas of the application. The application can write pixels into the canvas, and the display server will periodically read the canvas and write any changes to the LFB.

The display server determines the ordering of applications along the z-axis, drawing the pixels of the topmost application last. It also determines which application is currently active and directs key presses only to it.

It is highly important that the display server is efficient; if it runs too slowly, the system will have a low frame rate, with interaction feeling choppy. We had to carefully optimize the code for the display server before we got an acceptable frame rate inside QEMU.

We wrote a mouse driver for a PS/2 mouse which reads data packets from the mouse through PIO [14]. These data packets are generated whenever the mouse is moved. They only contain delta x- and y-coordinates, which means that the concept of a cursor only exists inside the display server.

8.5 SAMPLE GRAPHICAL APPLICATIONS

To test our design and implementation we have written two simple graphical applications: a paint program and a terminal emulator.

The paint program allows the user to click its canvas, drawing a square as a result.

The terminal emulator accepts keyboard input. It renders the input as the user types it. When the enter key is pressed, the terminal emulator forwards the input to a non-graphical process running the simple shell which we described in section 6.3. The output of this shell process is displayed to the user inside the terminal emulator.

An interaction with the two applications is displayed in figure **??**.

HARDWARE

In this lab we saught to run our operating system on real hardware rather than the QEMU emulator.

The main cause of difficulty was that our chosen hardware differed from the hardware emulated by QEMU. Additionally, there are aspects of a machine that QEMU does not emulate faithfully for efficiency reasons, but which matter when writing an operating system. This caused latent bugs to surface once we switched to real hardware.

In the following sections we describe the steps we took to get our operating system running on our machine of choice, a Packard Bell Dot S netbook [15].

9.1 BOOTING FROM USB

To boot our operating system we had to write its raw disk image to a bootable medium, such as a hard drive or a USB drive. We preferred to boot from a USB drive, since the hard drive would require either cable access or booting into a different operating system to transfer the image, which would be slow and tedious.

When we supplied our raw disk image to QEMU it booted as expected. We therefore expected that the same would happen when we transferred the image to a USB drive and inserted it into the netbook. However the USB was not recognized as a bootable medium.

The reason turns out to be historical [16]: when USB technology was new, there was disagreement between manifacturers on whether a USB drive should be a raw storage drive or a bootable medium. Initially the user could decide through a BIOS setting, but for usability reasons, manufacturers eventually opted to use heuristics to detect whether a USB is bootable or not.

These heuristics are based on data structures called the Master Boot Record (MBR) and BIOS Parameter Block (BPB). Both the MBR and BPB must be present in the first sector of a USB drive, otherwise most machines will assume that the drive is merely for data storage.

Once we added a valid MBR and BPB to our boot loader, the USB was recognized as bootable and the boot loader executed. Unfortunately it got stuck. The problem was that the boot loader was programmed to load the operating system from a hard drive connected through an ATA cable.

We realized that to continue, we would have to rewrite our boot loader to support loading from USB. However the boot loader is contrained to 512 bytes, since the BIOS loads it from the first sector of the drive. Writing code that can load from both USB and ATA drives in 512 bytes seemed difficult.

In the end we replaced our boot loader with a better one, namely GRUB, which supports booting from various hard drive types, USB, and even an ethernet connection.

We had to modify our file system to encorporate GRUB; the file system assumes that the boot loader will fit into the first 512 bytes, with the superblock residing in the following sector. Since GRUB uses multiple stages it needs much more space. We therefore had to modify our file system such that the first 32 MB are reserved for GRUB, and the superblock resides thereafter.

After integrating GRUB into the project we were finally able to boot directly from USB and have kernel code loaded. Unfortunately the scheduler seemed broken; it was unable to switch in new processes.

9.2 FINDING THE LAPIC

Debugging showed that timer interrupts were not being generated, which prevented the scheduler from preempting the running process. The LAPIC is used to generate these timer interrupts.

In section 5.4 we described how our kernel uses the so-called mpconfig method to find the LAPIC; that is, it looks for an MP configuration table in memory to find the physical address of the LAPIC. This address is used to communicate with the LAPIC, asking it to generate timer interrupts. However our kernel failed to find these configuration tables. They were simply not present in the memory of the netbook.

It turns out that the mpconfig method is outdated and unsupported by modern hardware. We therefore had to implement a different way to find the LAPIC.

We leave out the details of the method, but in essence newer machines contain so-called ACPI tables [17]. One of these is the APIC table, which contains the physical addres of the LAPIC. The ACPI tables can be found by scanning a region of memory for a data structure called the RSDP [18], which points at another data structure, the RSDT, which finally points at the ACPI tables.

After implementing this finding and parsing of ACPI tables, the kernel was able to discover the LAPIC and timer interrupts were generated, letting the scheduler work as intended.

9.3 LATENT BUGS

Our kernel was behaving oddly; modifying a PTE seemed to have no effect. After the modification, writing to memory at the virtual address still affected memory at the old physical address rather than the new one.

Since this seemed to be a caching issue, we realized that it was likely related to the TLB. Eventually we learned that the invlpg instruction must be used to invalidate a TLB entry after a PTE has been modified. Otherwise the outdated TLB entry will continue to be used until it is evicted from the cache.

Interestingly, this bug never surfaced while running the kernel in QEMU. We assume that QEMU does not faithfully emulate the TLB for performance reasons.[1]

After fixing the page table bug the machine successfully booted into a graphical interface. However the mouse behaved oddly, jumping around when moved.

The PS/2 mouse driver we wrote during the graphics lab was at fault; the driver uses PIO to read from the mouse. Before each inb instruction, it is necessary to wait for the mouse to signal that it has sent another byte of data. The driver did not do this. However in QEMU these waits were not necessary; thus this was another instance of QEMU not emulating hardware perfectly. We inserted the necessary delays in the mouse driver.

9.4 THE DEVELOPMENT PROCESS

The bugs and issues we have described seem rather obvious when explained in hindsight. But in practice, debugging a live system without a debugger was tricky and time-consuming. This was aggravated by there being a sizable delay between writing code and getting feedback from the machine, caused by us having to write the kernel image to a USB drive and switch it between machines for every code change.

For the page table bug, we only observed that processes behaved oddly when writing to memory. It took a lot of time and print statements before

---

[1] VirtualBox behaved similarly to QEMU; we therefore believe that this quirk can be used as a means to detect virtualization/emulation.

we realized the root cause of the problem. The other bugs and issues were similar.

In conclusion, there are significant advantages to developing an operating system on an emulator rather than a physical machine, despite emulators not being entirely accurate. Moving the operating system to real hardware was a time-consuming but not conceptually difficult process; the main task was supporting different hardware types and standards.

9.5 THE FINAL RESULT

With all the changes and fixes described so far, our operating system finally ran on the netbook and behaved as expected. A picture of the final system is shown in figure **??**.

CONCLUSION

Through writing an operating system from scratch we have gained a deep understanding of the many tasks that it must perform.

We have understood the boot process and the importance of a good boot loader. We have investigated the intricacies of physical page management and address translation, which were only a small part of the infrastructure needed to run a user space process. We have extended this infrastructure to support multiple interacting processes. We then wrote various user space programs and implemented a custom file system for them to reside on.

We then saw how various features, such as the network stack and the file system code, could be moved from the kernel to user space according to exokernel principles.

By designing and implementing a graphical user interface for our system we have additionally seen how graphics work, all the way from drawing a single pixel, to the end result with a display server managing windows and rendering fonts. As part of this process we gained experience writing drivers for various devices, including a networking card and a PS/2 mouse.

By moving our operating system to a real machine rather than an emulator we have seen the advantages emulators bring to the development process, as well as their limitations.

Throughout this process we have been repeatedly surprised by the scope of services that an operating system must provide, and which we previously took entirely for granted.

REFERENCES

[1] Frans Kaashoek and Robert Morris. 6.828 operating system engineering course. `https://pdos.csail.mit.edu/6.828/2016/index.html`.

[2] OSDev Wiki. Global descriptor table. `http://wiki.osdev.org/Global_Descriptor_Table`.

[3] OSDev Wiki. Cmos. `http://wiki.osdev.org/CMOS`.

[4] Intel Corporation. 80386 programmer's reference manual, chapter 5: Memory management. `https://pdos.csail.mit.edu/6.828/2016/readings/i386/c05.htm`.

[5] Intel Corporation. 80386 programmer's reference manual, chapter 9: Exceptions and interrupts. `https://pdos.csail.mit.edu/6.828/2016/readings/i386/c09.htm`.

[6] OSDev Wiki. Exceptions. `http://wiki.osdev.org/Exceptions`.

[7] Intel Corporation. *Intel 64 and IA-32 architectures software developer's manual*, volume 3A of *System programming guide*. October 2017.

[8] OSDev Wiki. Apic. `http://wiki.osdev.org/APIC`.

[9] Intel Corporation. *MultiProcessor Specification*. May 1997.

[10] Intel Corporation. *PCI/PCI-X Family of Gigabit Ethernet Controllers Software Developer's Manual*. March 2009.

[11] OSDev Wiki. Graphics stack. `http://wiki.osdev.org/Graphics_stack`.

[12] Windowing system. `https://en.wikipedia.org/wiki/Windowing_system`.

[13] X window system. `https://en.wikipedia.org/wiki/X_Window_System`.

[14] OSDev Wiki. Mouse input. `http://wiki.osdev.org/Mouse_Input`.

[15] Packard bell dot s. `http://www.packardbell.com/pb/en/IL/content/model/LU.BUNOD.038`.

[16] OSDev Wiki. Problems booting from usb flash. `http://wiki.osdev.org/Problems_Booting_From_USB_Flash`.

[17] OSDev Wiki. Symmetric multiprocessing. `wiki.osdev.org/Symmetric_Multiprocessing`.

[18] OSDev Wiki. Rsdp. `http://wiki.osdev.org/RSDP`.